

Estructura del Libro

Official Google Cloud Certified Professional Machine Learning Engineer Study Guide

Índice de Contenidos

Cover

Table of Contents

List of Tables

List of Illustrations

Guide

Pages

**Official Google Cloud Certified Professional Machine
Learning Engineer Study Guide**

F02

F03

Acknowledgments

About the Author

About the Technical Editors

About the Technical Proofreader

Google Technical Reviewer

Introduction

Google Cloud Professional Machine Learning Engineer Certification

Why Become Professional ML Engineer (PMLE) Certified?

How to Become Certified

Who Should Buy This Book

How This Book Is Organized

Chapter Features

Bonus Digital Contents

Interactive Online Learning Environment and Test Bank

Conventions Used in This Book

Google Cloud Professional ML Engineer Objective Map

How to Contact the Publisher

Assessment Test

Answers to Assessment Test

Chapter 1 Framing ML Problems

Translating Business Use Cases

Machine Learning Approaches

Supervised, Unsupervised, and Semi-supervised Learning

Classification, Regression, Forecasting, and Clustering

ML Success Metrics

Area Under the Curve Receiver Operating Characteristic (AUC ROC)

The Area Under the Precision■Recall (AUC PR) Curve

Regression

Responsible AI Practices

Summary

Exam Essentials

Review Questions

Chapter 2Exploring Data and Building Data Pipelines

Visualization

Box Plot

Line Plot

Bar Plot

Scatterplot

Statistics Fundamentals

Mean

Median

Mode

Outlier Detection

Standard Deviation

Correlation

Data Quality and Reliability

Data Skew

Data Cleaning

Scaling

Log Scaling

Z-score

Clipping

Handling Outliers

Establishing Data Constraints

Exploration and Validation at Big Data Scale

Running TFDV on Google Cloud Platform

Organizing and Optimizing Training Datasets

Imbalanced Data

Data Splitting

Data Splitting Strategy for Online Systems

Handling Missing Data

Data Leakage

Summary

Exam Essentials

Review Questions

Chapter 3 Feature Engineering

Consistent Data Preprocessing

Encoding Structured Data Types

Why Transform Categorical Data?

Mapping Numeric Values

Normalizing

Bucketing

Mapping Categorical Values

Label Encoding or Integer Encoding

One-Hot Encoding

Out of Vocab (OOV)

Feature Hashing

Hybrid of Hashing and Vocabulary

Embedding

Feature Selection

Class Imbalance

Classification Threshold with Precision and Recall

Area under the Curve (AUC)

AUC ROC

AUC PR

Feature Crosses

TensorFlow Transform

TensorFlow Data API (tf.data)

TensorFlow Transform

GCP Data and ETL Tools

Summary

Exam Essentials

Review Questions

Chapter 4Choosing the Right ML Infrastructure

Pretrained vs. AutoML vs. Custom Models

Pretrained Models

Vision AI

Video AI

Natural Language AI

Translation AI

Speech■to■Text

Text■to■Speech

AutoML

AutoML for Tables or Structured Data

AutoML for Images and Video

AutoML for Text

Recommendations AI/Retail AI

Document AI

Dialogflow and Contact Center AI

Virtual Agent (Dialogflow)

Agent Assist

Insights

CCAI

Custom Training

How a CPU Works

GPU

TPU

How to Use TPUs

Advantages of TPUs

When to Use CPUs, GPUs, and TPUs

Cloud TPU Programming Model

Provisioning for Predictions

Scaling Behavior

Finding the Ideal Machine Type

Edge TPU

Deploy to Android or iOS Device

Summary

Exam Essentials

Review Questions

Chapter 5 Architecting ML Solutions

Designing Reliable, Scalable, and Highly Available ML Solutions

Choosing an Appropriate ML Service

Data Collection and Data Management

Google Cloud Storage (GCS)

BigQuery

Vertex AI Managed Datasets

Vertex AI Feature Store

NoSQL Data Store

Automation and Orchestration

Use Vertex AI Pipelines to Orchestrate the ML Workflow

Use Kubeflow Pipelines for Flexible Pipeline Construction

Use TensorFlow Extended SDK to Leverage Prebuilt Components for Common Steps

When to Use Which Pipeline

Serving

Offline or Batch Prediction

Online Prediction

Summary

Exam Essentials

Review Questions

Chapter 6 Building Secure ML Pipelines

Building Secure ML Systems

Encryption at Rest

Encryption in Transit

Encryption in Use

Identity and Access Management

IAM Permissions for Vertex AI Workbench

Securing a Network with Vertex AI

Securing Vertex AI Workbench

Securing Vertex AI Endpoints

Securing Vertex AI Training Jobs

Federated Learning

Differential Privacy

Format■Preserving Encryption and Tokenization

Privacy Implications of Data Usage and Collection

Google Cloud Data Loss Prevention

Google Cloud Healthcare API for PHI Identification

Best Practices for Removing Sensitive Data

Summary

Exam Essentials

Review Questions

Chapter 7Model Building

Choice of Framework and Model Parallelism

Data Parallelism

Synchronous Training

Asynchronous Training

Model Parallelism

Modeling Techniques

Artificial Neural Network

Deep Neural Network (DNN)

Convolutional Neural Network

Recurrent Neural Network

What Loss Function to Use

Gradient Descent

Learning Rate

Batch

Batch Size

Epoch

Hyperparameters

Tuning Batch Size

Tuning Learning Rate

Transfer Learning

Semi■supervised Learning

When You Need Semi■supervised Learning

Limitations of SSL

Data Augmentation

Offline Augmentation

Online Augmentation

Model Generalization and Strategies to Handle Overfitting and Underfitting

Bias Variance Trade■Off

Underfitting

Overfitting

Regularization

Summary

Exam Essentials

Review Questions

Chapter 8 Model Training and Hyperparameter Tuning

Ingestion of Various File Types into Training

Collect

Process

Cloud Dataflow

Cloud Data Fusion

Cloud Dataproc

Cloud Composer

Cloud Dataprep

Summary of Processing Tools

Store and Analyze

Developing Models in Vertex AI Workbench by Using Common Frameworks

Creating a Managed Notebook

Exploring Managed JupyterLab Features

Data Integration

BigQuery Integration

Ability to Scale the Compute Up or Down

Git Integration for Team Collaboration

Schedule or Execute a Notebook Code

Creating a User-Managed Notebook

Training a Model as a Job in Different Environments

Training Workflow with Vertex AI

Training Dataset Options in Vertex AI

Pre-built Containers

Custom Containers

Distributed Training

Hyperparameter Tuning

Why Hyperparameters Are Important

Techniques to Speed Up Hyperparameter Optimization

How Vertex AI Hyperparameter Tuning Works

Vertex AI Vizier

How Vertex AI Vizier Differs from Custom Training

Tracking Metrics During Training

Interactive Shell

TensorFlow Profiler

What-If Tool

Retraining/Redeployment Evaluation

Data Drift

Concept Drift

When Should a Model Be Retrained?

Unit Testing for Model Training and Serving

Testing for Updates in API Calls

Testing for Algorithmic Correctness

Summary

Exam Essentials

Review Questions

Chapter 9 Model Explainability on Vertex AI

Model Explainability on Vertex AI

Explainable AI

Interpretability and Explainability

Feature Importance

Vertex Explainable AI

Feature Attribution

Vertex AI Example-Based Explanations

Data Bias and Fairness

ML Solution Readiness

How to Set Up Explanations in the Vertex AI

Summary

Exam Essentials

Review Questions

Chapter 10 **Scaling Models in Production**

Scaling Prediction Service

TensorFlow Serving

Serving a Saved Model with TensorFlow Serving

Serving (Online, Batch, and Caching)

Real-time Static and Dynamic Reference Features

Precomputing and Caching Prediction

Google Cloud Serving Options

Online Predictions

Deploying the Model

Deploying a Model Using an API

Deploying a Model Using the Google Cloud Console

Make Predictions

A/B Testing of Different Versions of a Model

Undeploy Endpoints

[Send an Online Explanation Request](#)

Batch Predictions

Hosting Third-Party Pipelines (MLflow) on Google Cloud

Testing for Target Performance

Configuring Triggers and Pipeline Schedules

Summary

Exam Essentials

Review Questions

Chapter 11 Designing ML Training Pipelines

Orchestration Frameworks

Kubeflow Pipelines

Vertex AI Pipelines

Apache Airflow

Cloud Composer

Comparison of Tools

Identification of Components, Parameters, Triggers, and Compute Needs

Schedule the Workflows with Kubeflow Pipelines

Schedule Vertex AI Pipelines

System Design with Kubeflow/TFX

System Design with Kubeflow DSL

Kubeflow Pipelines Components

System Design with TFX

Hybrid or Multicloud Strategies

Summary

Exam Essentials

Review Questions

Chapter 12 Model Monitoring, Tracking, and Auditing Metadata

Model Monitoring

Concept Drift

Data Drift

Model Monitoring on Vertex AI

Drift and Skew Calculation

Practical Considerations of Enabling Monitoring

Input Schemas

Automatic Schema Parsing

Custom Schema

Logging Strategy

Types of Prediction Logs

Container Logging

Access Logging

Request■Response Logging

Log Settings

Model Monitoring and Logging

Model and Dataset Lineage

Vertex ML Metadata

Manage ML Metadataschemas

Vertex AI Pipelines

Vertex AI Experiments

Vertex AI Debugging

Summary

Exam Essentials

Review Questions

Chapter 13Maintaining ML Solutions

MLOps Maturity

MLOps Level 0: Manual/Tactical Phase

Key Features of Level 0

Challenges

MLOps Level 1: Strategic Automation Phase

Key Features of MLOps Level 1

Challenges

MLOps Level 2: CI/CD Automation, Transformational Phase

Key Features of Level 2

Retraining and Versioning Models

Triggers for Retraining

Versioning Models

Feature Store

Solution

Data Model

Ingestion and Serving

Vertex AI Permissions Model

Custom Service Account

Access Transparency in Vertex AI

Common Training and Serving Errors

Training Time Errors

Serving Time Errors

TensorFlow Data Validation

Vertex AI Debugging Shell

Summary

Exam Essentials

Review Questions

Chapter 14BigQuery ML

BigQuery – Data Access

BigQuery ML Algorithms

Model Training

Model Evaluation

Prediction

Explainability in BigQuery ML

BigQuery ML vs. Vertex AI Tables

Interoperability with Vertex AI

Access BigQuery Public Dataset

Import BigQuery Data into Vertex AI

Access BigQuery Data from Vertex AI Workbench Notebooks

Analyze Test Prediction Data in BigQuery

Export Vertex AI Batch Prediction Results

Export BigQuery Models into Vertex AI

BigQuery Design Patterns

Hashed Feature

Transforms

Summary

Exam Essentials

Review Questions

Appendix Answers to Review Questions

Chapter 1: Framing ML Problems

Chapter 2: Exploring Data and Building Data Pipelines

Chapter 3: Feature Engineering

Chapter 4: Choosing the Right ML Infrastructure

Chapter 5: Architecting ML Solutions

Chapter 6: Building Secure ML Pipelines

Chapter 7: Model Building

Chapter 8: Model Training and Hyperparameter Tuning

Chapter 9: Model Explainability on Vertex AI

Chapter 10: Scaling Models in Production

Chapter 11: Designing ML Training Pipelines

Chapter 12: Model Monitoring, Tracking, and Auditing Metadata

Chapter 13: Maintaining ML Solutions

Chapter 14: BigQuery ML

Index

A

B

C

D

E

F

G

H

I

J

K

L

M

N

O

P

Q

R

S

T

U

V

W

Z

Online Test Bank

Register and Access the Online Test Bank

WILEY END USER LICENSE AGREEMENT

Total de capítulos encontrados: 433

Archivo EPUB: my_epub.epub